

# The Shape of Learning in Repeated Games

Information-Geometric Analysis via InPCA

Nora Han

Supervised by Prof. Aaron Roth

Learning in Games

CIS 6200 Advanced Topics in Machine Learning

Dec 4, 2025



# Outline

Introduction

Background and Prior Work

Experimental Setup

Visualization Pipeline

Results

MWU vs Constant — Curve

OGD vs OGD — Hexagonal Cycle

OGD vs MWU — Triangular Cycle

Probability and Geomtry

Discussion and Future Work

Reference

# Intrduction

- 1.1 Problem Setting
- 1.2 Why Learning Trajectories Matter

## 1.1 Problem Setting

- Repeated-play games: agents update strategies over time.
- Standard analysis tracks regret and convergence, not *how* strategies move.
- Learning algorithms (MWU, OGD, etc.) produce complex trajectories in probability space.
- Existing projection (naive PCA) distort or hide this structure.

## 1.2 Why Learning Trajectories Matter

- Last-iterate may cycle even when time-average converges.
- Trajectory geometry reveals exploitability, stability, and dynamic modes.
- Different algorithms trace distinct shapes—this is not captured by regret alone.
- Understanding geometry => understanding algorithmic behavior.

### Core Question

*What hidden geometric structure underlies learning in repeated games?*

# Background and Prior Work

- 2.1 Classical Results on Convergence in Repeated Games
- 2.2 No-Regret Learning Rules and Known Behaviors
- 2.3 Limitations of Existing Visualization and Analysis Approaches

## 2.1 Classical Results on Convergence in Repeated Games

### Repeated zero-sum game setup

- Two players (row, column), finite action sets  $A_1, A_2$  with  $|A_1| = m, |A_2| = n$ .
- Mixed strategies  $p_t \in \Delta_m, q_t \in \Delta_n$ ; payoff matrix  $A \in \mathbb{R}^{m \times n}$ .
- Stage- $t$  payoff to row:

$$u(p_t, q_t) = p_t^\top A q_t.$$

## Last-iterate vs time-average

- *Last iterate*: the actual strategy profile at time  $T$ ,

$$(p_T, q_T).$$

- *Time-average (empirical distribution)*:

$$\bar{p}_T := \frac{1}{T} \sum_{t=1}^T p_t, \quad \bar{q}_T := \frac{1}{T} \sum_{t=1}^T q_t.$$

- We say *time-average converges* if  $(\bar{p}_T, \bar{q}_T)$  approaches the minimax set as  $T \rightarrow \infty$ .
- We say *last iterate converges* if  $(p_T, q_T)$  converges to a limit point (e.g., a Nash equilibrium).

# Classical guarantees

- Minimax theorem:

$$\max_{p \in \Delta_m} \min_{q \in \Delta_n} p^\top A q = \min_{q \in \Delta_n} \max_{p \in \Delta_m} p^\top A q =: v^*.$$

- If both players use no-regret algorithms, the empirical play distribution converges to the set of coarse correlated equilibria; in zero-sum games,  $(\bar{p}_T, \bar{q}_T)$  approaches the minimax set and achieves near-optimal value  $v^*$ .
- However, recent work on FTRL (Follow-the-Regularized-Leader) dynamics shows a sharp separation:
  - Time-averages converge to equilibrium-like objects.
  - The *actual trajectory*  $(p_t, q_t)$  can be *recurrent* (Poincaré cycles) and fail to converge.

# Takeaway

- Many learning algorithms **look like they converge** when we measure averages over time.
- But if we watch what the algorithms are actually *doing each step*, their strategies often **move** rather than settle down.
- These geometric patterns are **not visible** through regret bounds or equilibrium guarantees.

## Why This Matters

To truly understand learning in games, we look at *the shape of the path*, not just where the averages end up.

## 2.2 No-Regret Learning Rules and Known Behaviors

### Multiplicative Weights

- Action set  $A = \{1, \dots, k\}$ . Maintain weights  $w_{t,i} > 0$  and mixed strategy

$$p_t(i) = \frac{w_{t,i}}{\sum_{j=1}^k w_{t,j}}.$$

- At round  $t$ , observe a cost (or loss) vector  $c_t \in \mathbb{R}^k$ .
- Update rule, where  $\eta > 0$  is a learning rate:

$$w_{t+1,i} = w_{t,i} \exp(-\eta c_{t,i}), \quad p_{t+1}(i) \propto p_t(i) \exp(-\eta c_{t,i}),$$

## Intuition

- Treat each action like a candidate you are voting for.
- Actions that perform poorly **lose support exponentially fast**.
- Actions that perform well **keep or gain probability**.
- You never fully commit — you keep a soft distribution and shift attention toward winners.
- It is like *putting your money on what worked*, but gently.

## Core Idea

MWU is an exponential reward-shifting process. It learns by *rewarding success and punishing failure multiplicatively*.

## Online Gradient Descent (OGD)

- Actions  $a_t \in K \subset \mathbb{R}^d$  (convex). At round  $t$ , observe cost vector  $c_t$  with  $\|c_t\|_2 \leq C$ .
- Unconstrained update:

$$a_{t+1} = a_t - \frac{\eta}{2} c_t.$$

- With constraints, project back to  $K$ :

$$a_{t+1} = \Pi_K \left( a_t - \frac{\eta}{2} c_t \right).$$

## Intuition

- Think of the strategy as a point on a landscape.
- Each round tells you the *direction of steepest improvement*.
- OGD takes a **straight step** downhill in that direction.
- If the step leaves the allowed region (the simplex), we **snap back** to the closest valid point.
- It is like *nudging your strategy a tiny bit toward something better*, every step.

## Core Idea

OGD learns by *moving linearly in the direction that reduces loss the fastest*.

## Known behaviors in games

- If both players run no-regret algorithms (MWU, OGD, FTRL variants), then:
  - Empirical play (time-average) converges to equilibrium-like sets (Nash/CCE) in zero-sum and more general games.
  - But the *last iterate* often *cycles* around interior equilibria.
- These dynamics can be fast-regret and yet non-convergent in last iterate — a key motivation for studying their geometric trajectories.

## 2.3 Limitations of Existing Visualization and Analysis Approaches

### Regret-based analysis as a “black box”

- Standard theory: if algorithms are no-regret, empirical play converges to CCE / Nash-like objects.
- This guarantees *what averages do*, but not *how the last iterate moves* in state space.
- Convergent, recurrent, and even chaotic systems can all have similarly low regret.

## Phase portraits and payoff traces

- Continuous-time dynamics (replicator, FTRL ODEs) can be visualized via 2D phase portraits in simple  $2 \times 2$  or symmetric games.
- In higher-dimensional, discrete-time repeated games (e.g.,  $3 \times 3$  Shapley), phase portraits become:
  - High-dimensional and hard to project meaningfully.
  - Sensitive to choice of coordinates and projections.
- Payoff vs. time plots only show *one scalar* per round; they lose structural information about the shape of the trajectory.

# Why Intensive Principal Component Analysis(InPCA) for Probability Dynamics

**The challenge** - Strategies are probability distributions. Their natural geometry is curved, not Euclidean.

- Mixed strategies live in a probability simplex

$$\Delta^{k-1} = \{p \in \mathbb{R}^k : p_i \geq 0, \sum_i p_i = 1\},$$

which is not a flat space.

- Changing one probability forces another to adjust — directions are not independent.
- Natural distances between strategies are information-theoretic (KL, Hellinger, Fisher metric), not Euclidean.
- These distances induce a **curved geometric manifold** structure.

## Implication

Learning trajectories live on a curved surface. Flattening them with PCA distorts their shape.

## What goes wrong with PCA

- PCA assumes data lives in a flat space with straight-line distances.
- Probability vectors live on a **simplex**, where distances behave like KL/Hellinger divergences, not Euclidean norms.
- When we apply PCA to strategy trajectories, it **distorts** the geometry: cycles smear into blobs; polygons collapse into clouds.

## What InPCA does

- Uses a kernel derived from the information geometry of distributions.
- Respects **intrinsic distances** between probability states.
- Embeds trajectories into a space where **learning updates appear as clean geometric shapes**.

# Experimental Setup

- 3.1 Game Environments
- 3.2 Learning Algorithms
- 3.3 Logging Joint Action Distributions

## 3.1 Game Environments

### Rock-Paper-Scissors (Zero-Sum $3 \times 3$ )

- Classical benchmark for cycling behavior.
- Unique mixed Nash equilibrium at  $(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$ .
- Simple rules, surprisingly rich dynamics.

### Shapley's Non-Transitive Game ( $3 \times 3$ )

- General-sum payoff matrix with no static dominance ordering.
- Known to generate *limit cycles* under standard learning rules.
- Serves as a stress test for convergence claims.

## Shapley's 3×3 Game

**Shapley's Game** (Shapley, 1964) is a  $3 \times 3$  general-sum game where each player's best response cycles. No action is globally optimal.

$$A_{\text{row}} = \begin{pmatrix} 0 & 1 & -1 \\ -1 & 0 & 1 \\ 1 & -1 & 0 \end{pmatrix}, \quad A_{\text{col}} = -A_{\text{row}} + \epsilon I$$

Each action is beaten by exactly one other action.

- There exists a unique mixed Nash equilibrium, but **best responses chase each other in cycles**.
- No-regret learning does *not* stabilize — instead, strategies trace persistent orbits.

## 3.2 Learning Algorithms

### Multiplicative Weights Update (MWU)

- Increases probability of actions that performed well.
- Grows and shrinks weights exponentially based on payoff feedback.
- Dynamics tend to be smooth and rotational.

### Online Gradient Descent (OGD)

- Moves strategies in the direction that reduces loss.
- Then projects back to the probability simplex.
- Can produce abrupt directional changes due to projection—leading to corner-like turns.

### Constant Strategies (Baseline)

- One player never updates their distribution.
- Used to reveal pure-response structure of the adaptive learner.

### 3.3 Logging Joint Action Distributions

- At each round, we record the **joint mixed strategy**

$$P_t = p_t \otimes q_t \in \mathbb{R}^9,$$

representing the full distribution over action pairs.

- This lifts the analysis from individual beliefs to their interaction space.
- Each  $P_t$  lies on the probability simplex—a curved manifold.
- These logged vectors become the raw material for our geometric embeddings.

**Why Joint Distributions?** The game is not defined by  $p_t$  or  $q_t$  alone, but by their *coupled evolution*. Geometry emerges in the joint space.

# Visualization Pipeline

- 4.1 Joint Play as a 9-D Probability Vector
- 4.2 Embedding  $P_t$  Using InPCA
- 4.3 3D Trajectories Over Time

## 4.1 Joint Play as a 9-D Probability Vector

**Goal:** Represent what the two players *jointly* do at each timestep.

- Each player chooses a mixed strategy:

$$p_t \in \Delta^2 \quad q_t \in \Delta^2,$$

where  $\Delta^2$  is the 3-action simplex (R, P, S).

- The joint play is a probability distribution over action *pairs*:

$$P_t = p_t \otimes q_t \in \Delta^8 \quad (\text{vector of length } 3 \times 3 = 9).$$

- Entry  $P_t(i, j)$  = probability row plays action  $i$  *and* column plays action  $j$ .

## Why 9-D?

Pairwise interactions—not individual strategies—determine payoffs, regret, and exploitability. The *outer product* preserves all second-order structure that is invisible in marginal plots of  $p_t$  or  $q_t$  alone.

## 4.2 Embedding $P_t$ Using InPCA

**Input:** Joint distributions  $P_t \in \Delta^8$ .

1. Hellinger lift:

$$S_t = \sqrt{P_t}$$

2. Center the cloud:

$$\tilde{S}_t = S_t - \frac{1}{T} \sum_{s=1}^T S_s$$

3. SVD / PCA:

$$\tilde{S} = U \Sigma V^\top, \quad X_t = U_{t,1:3} \cdot \frac{\Sigma_{1:3}}{\sqrt{2}}$$

# How InPCA Finds 3D Geometry

**Goal:** Reduce a 9-dimensional probability trajectory to 3 meaningful geometric axes.

## 1. Collect all centered points

$$\tilde{S} \in \mathbb{R}^{T \times 9} = \begin{bmatrix} \tilde{S}_1^\top \\ \tilde{S}_2^\top \\ \vdots \\ \tilde{S}_T^\top \end{bmatrix}$$

Each row is the square-rooted, mean-centered joint distribution at time  $t$ .

## 2. Find dominant directions of variation (SVD)

$$\tilde{S} = U\Sigma V^\top$$

- $V$  contains *principal directions* in the 9-D space.
- $\Sigma$  tells us how important each direction is.
- $U$  gives each timestep's coordinates along those directions.

## 3. Keep only the top 3 (the “geometric modes”)

$$x_t = U_{t,1:3} \cdot \frac{\Sigma_{1:3}}{\sqrt{2}} \in \mathbb{R}^3$$

### What this means

We are not plotting raw probabilities. We are plotting the *three strongest ways the strategy profile changes over time*. These axes reveal rings, polygons, and turning points that ordinary PCA and regret curves cannot see.

## Interpretation

InPCA embeds probability dynamics into a Euclidean space that respects statistical distance. The resulting 3D coordinates reveal the geometric structure of learning trajectories.

## Why InPCA?

Standard PCA treats probabilities like flat vectors. InPCA respects statistical distance, revealing geometric structure that PCA *hides*.

# 3D Trajectories Over Time

**Coordinates:**  $(x_t^{(1)}, x_t^{(2)}, x_t^{(3)})$  for each timestep  $t$ .

- Normalize time:

$$\tau = \frac{t - t_{\min}}{t_{\max} - t_{\min}} \in [0, 1].$$

- Plot trajectories of (PC1, PC2, PC3) in 3D.
- Compare “last iterate” vs. “time-average” as separate layers.

## Interpretation

The radius, direction, and turns of the trajectory expose *how* algorithms learn—not just *where* they end. Shapes become directly visible.

# Why Normalize Time?

- Different runs take the same number of iterations, but *their internal progress is not synchronized*. Early steps explore rapidly; later ones move slowly.
- Raw timestep  $t$  is meaningless across runs:

$t = 150$  may be “early” in one run and “late” in another.

- Normalize time to a common scale:

$$\tau = \frac{t - t_{\min}}{t_{\max} - t_{\min}} \in [0, 1]$$

so that:

- $\tau = 0$  always means “beginning of learning”,
- $\tau = 1$  always means “steady state / late behavior”,
- trajectories from different runs become directly comparable.
- This alignment allows us to compute meaningful averages, overlay replicas, and reveal shared geometric patterns.

# Results

- 5.1 MWU vs Constant — Converge Geometry
- 5.2 OGD vs OGD — Hexagonal Cycle
- 5.3 OGD vs MWU — Triangular Cycle
- 5.4 Probability and Geometry

# 5.1 MWU vs Constant – Converge Geometry

- Setup: row uses MWU, column plays a fixed pure strategy (e.g. always Scissors).
- The joint distribution  $P_t = p_t \otimes q_{\text{const}}$  explores a *curved line* in InPCA space.

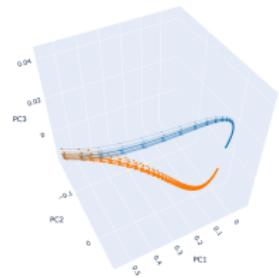


Figure 1: InPCA trajectory for MWU (row) vs constant (col)

## 5.1 MWU vs Constant – Radius over Time

- Radius in InPCA space:

$$\text{radius}_t = \sqrt{\text{PC}_1^2 + \text{PC}_2^2 + \text{PC}_3^2}$$

- Last-iterate radius remains roughly stable  $\Rightarrow$  persistent cycling.
- Time-average radius decays toward the center  $\Rightarrow$  approach to equilibrium in average play.

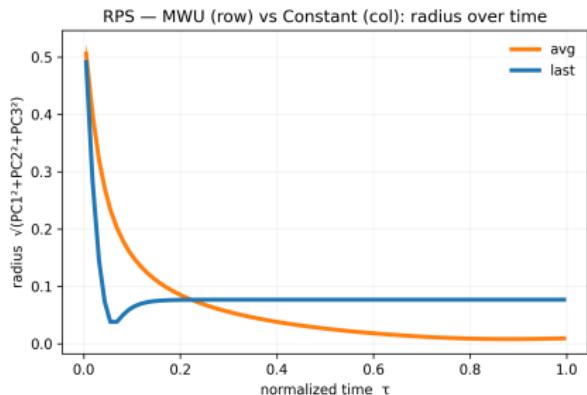


Figure 2: Radius vs. normalized time  $\tau$  for MWU vs constant.

## 5.2 OGD vs OGD – Hexagonal Cycle

- Both players run OGD in zero-sum RPS.
- InPCA reveals a *hexagon-like* orbit in the joint distribution  $P_t$ .
- Vertices of the hexagon correspond to moments when one player is nearly pure in Rock/Paper/Scissors.
- The path is more “cornered” and piecewise-linear than MWU’s smooth ring, reflecting OGD’s projection step.

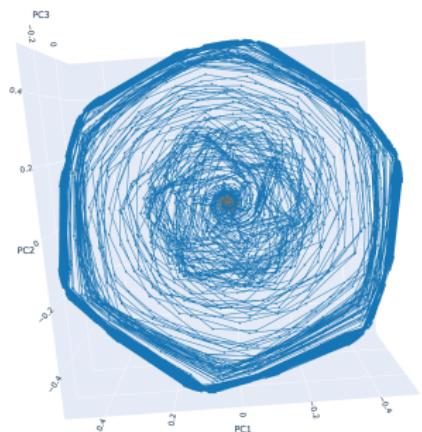


Figure 3: OGD (row) vs OGD (col): hexagon

## 5.2 OGD vs OGD – Radius and Angular Structure

- Radius over time stays roughly constant for the last iterate  $\Rightarrow$  stable limit cycle.
- Angle histogram on the  $(PC_1, PC_2)$  plane shows a strong  $k \approx 6$  harmonic.
- This quantitatively matches the visual hexagon: six preferred directions in the orbit.

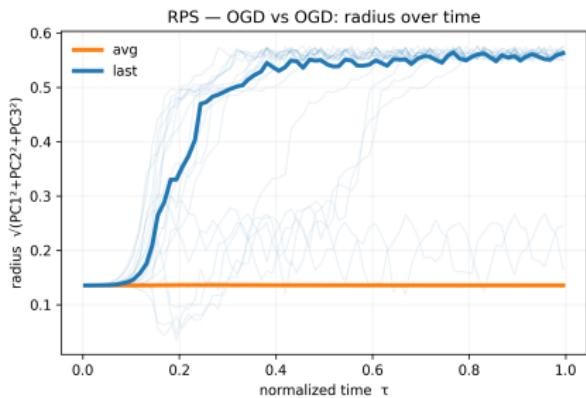


Figure 4: radius vs. time for OGD vs OGD

## 5.3 OGD vs MWU – Triangular Cycle

- Row uses OGD, column uses MWU on RPS.
- InPCA trajectory now concentrates on a *triangle-like* cycle instead of a ring or hexagon.
- Each corner corresponds to OGD's strategy being nearly pure (Rock, Paper, or Scissors).
- MWU's smoother dynamics + OGD's projected updates combine into a three-cornered orbit.

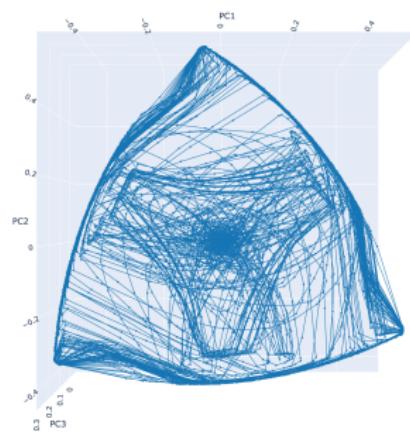


Figure 5: OGD (row) vs MWU (col): triangle

## 5.3 OGD vs MWU — Radius and Fourier Modes

- Radius for the last iterate again remains roughly constant: a robust limit cycle.
- Angle histogram in the  $(PC_1, PC_2)$  plane shows a dominant  $k \approx 3$  mode.
- Comparing  $k \approx 1$  (curve),  $k \approx 6$  (hexagon), and  $k \approx 3$  (triangle) quantifies the geometry.

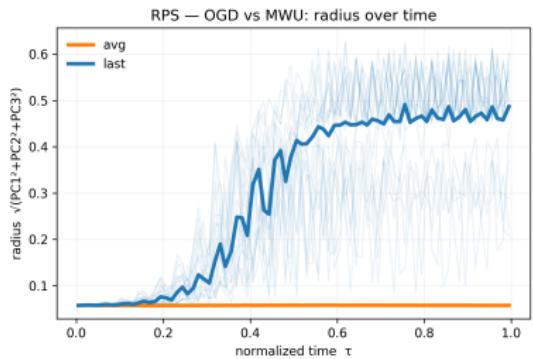


Figure 6: radius vs. time for OGD vs MWU

## 5.4 Probability and Geometry

- We plot out probabilities, and look at the probabilities at the turning points.
- For OGD, these turning points lie near the *pure strategies*:  
 $(1, 0, 0), (0, 1, 0), (0, 0, 1)$ .
- OGD repeatedly drives strategies to the simplex vertices.

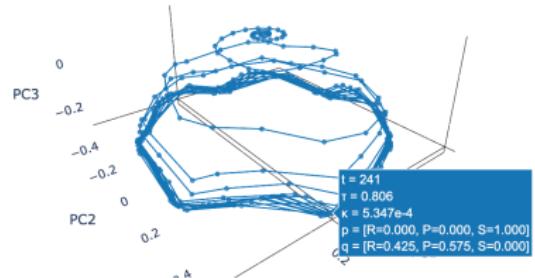


Figure 7: Curvature maxima align with pure actions under OGD

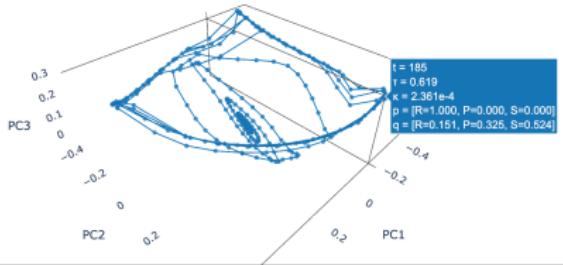


Figure 8: Curvature maxima align with pure actions under OGD for OGD v.s. MWU

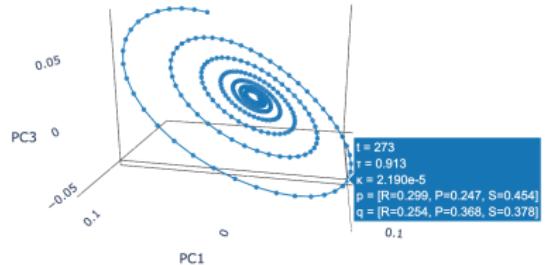


Figure 9: Reference for MWU vs MWU

# Why PCA Cannot Reveal Polygonal Geometry

- Standard PCA treats joint distributions  $P_t$  as Euclidean vectors. It ignores the **information geometry** of probabilities: the simplex is curved, and distances are not linear.
- Variance along  $\sqrt{P_t}$  directions collapses under PCA, flattening distinct modes of motion.
- Thus, PCA answers "*where the mass is*" but not "*how the policy moves*."

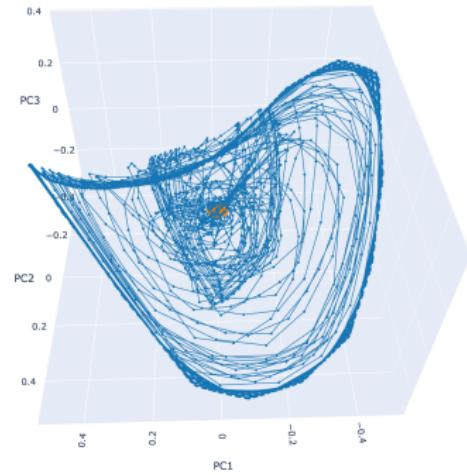


Figure 10: PCA collapses cycles into clouds for OGD vs MWU.

# Discussion and Future Work

- **Hypothesis: Why OGD Creates Edges**

- OGD performs *linear* gradient steps, then projects back onto the probability simplex.
- Projection frequently snaps strategies toward simplex *faces and vertices*, producing *piecewise-linear trajectories*.
- MWU updates are *multiplicative and smooth*, avoiding hard projections—hence no sharp turns or polygonal edges.

- **Beyond Pairwise Learning**

- Current pipeline tracks  $(p_t, q_t)$  and their outer-product  $P_t$ .
- Future work: log cross-derivatives, regret dynamics, and correlations to uncover *higher-order structure* in learning flows.

- **Algorithmic Extensions**

- Test other no-regret rules: Follow-The-Regularized-Leader, Optimistic OGD, Replicator Dynamics.
- Does each algorithm leave a unique “geometric fingerprint” in InPCA space?

# References

-  K. N. Quinn, C. B. Clement, F. De Bernardis, M. D. Niemack, & J. P. Sethna,  
*Visualizing probabilistic models and data with Intensive Principal Component Analysis*,  
Proceedings of the National Academy of Sciences (PNAS), 116(28):13762–13767, 2019.  
doi: [10.1073/pnas.1817218116](https://doi.org/10.1073/pnas.1817218116).
-  P. Mertikopoulos, C. Papadimitriou, & G. Piliouras,  
*Cycles in Adversarial Regularized Learning*,  
In *Proceedings of the 2018 ACM-SIAM Symposium on Discrete Algorithms (SODA)*,  
pp. 2703–2717, SIAM, 2018.  
Also available as arXiv:1709.02738.
-  A. Roth,  
*Learning in Games (and Games in Learning)*, Draft textbook, University of Pennsylvania,  
2024. provided as course material  
Provides definitions and analysis of MWU, OGD, and regret-based dynamics.